

EXPERIMENT:**Install, Configure and Run Hadoop and HDFS****PROGRAM:**

AIM: To Installing and Running Applications On Hadoop and HDFS.

HADOOP INSTALATION IN WINDOWS**1. Prerequisites**

Hardware Requirement

* RAM — Min. 8GB, if you have SSD in your system then 4GB RAM would also work.

* CPU — Min. Quad core, with at least 1.80GHz

2. JRE 1.8 — Offline installer for JRE

3. Java Development Kit — 1.8

4. A Software for Un-Zipping like 7Zip or Win Rar

* I will be using a 64-bit windows for the process, please check and download the version supported by your system x86 or x64 for all the software.

5. Download Hadoop zip

* I am using Hadoop-2.9.2, you can use any other STABLE version for hadoop.

Index of /dist/hadoop/core/hadoop-2.9.2

Name	Last modified	Size	Description
Parent Directory	-	-	-
hadoop-2.9.2-src.tar.gz	2020-07-03 04:37	37M	
hadoop-2.9.2-src.tar.gz.asc	2020-07-03 04:36	801	
hadoop-2.9.2-src.tar.gz.mds	2020-07-03 04:36	1.0K	
hadoop-2.9.2.tar.gz	2020-07-03 04:38	349M	
hadoop-2.9.2.tar.gz.asc	2020-07-03 04:37	801	
hadoop-2.9.2.tar.gz.mds	2020-07-03 04:36	1.0K	

Fig. 1:- Download Hadoop 2.9.2

Once we have Downloaded all the above software, we can proceed with next steps in installing the Hadoop.

2. Unzip and Install Hadoop

After Downloading the Hadoop, we need to Unzip the hadoop-2.9.2.tar.gz file.

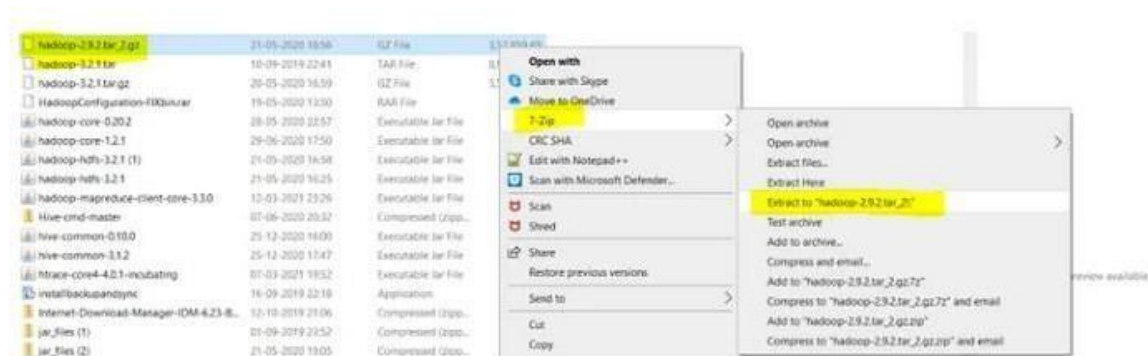


Fig. 2:- Extracting Hadoop Step-1

Once extracted, we would get a new file `hadoop-2.9.2.tar`.

Now, once again we need to extract this tar file.




Fig. 3:- Extracting Hadoop Step-2

Now we can organize our Hadoop installation, we can create a folder and move the final extracted file in it. For Eg. :-

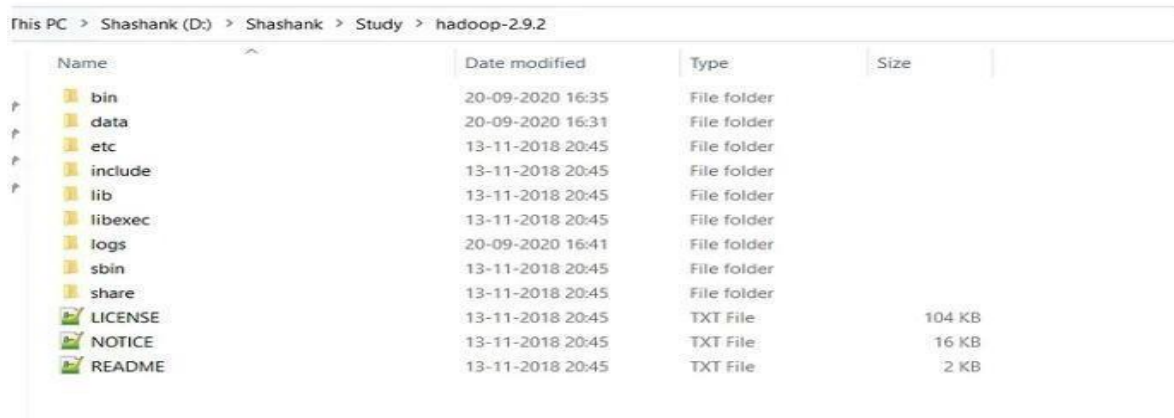


Fig. 4:- Hadoop Directory

Please note while creating folders, DO NOT ADD SPACES IN BETWEEN THE FOLDER NAME.(it can cause issues later)

I have placed my Hadoop in D: drive you can use C: or any other drive also.

3. Setting Up Environment Variables

Another important step in setting up a work environment is to set your Systems environment variable.

To edit environment variables, go to Control Panel > System > click on the -Advanced system settings link

Alternatively, We can Right click on This PC icon and click on Properties and click on the -Advanced system settings link

Or, easiest way is to search for Environment Variable in search bar and there you GO...

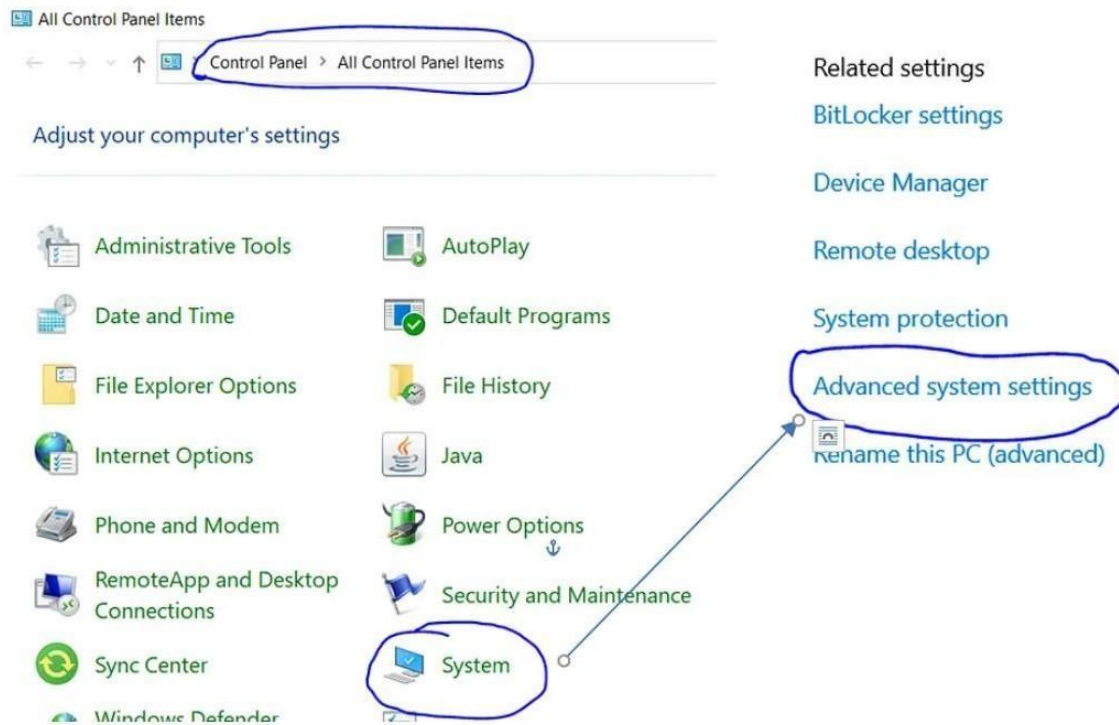


Fig. 5:- Path for Environment Variable

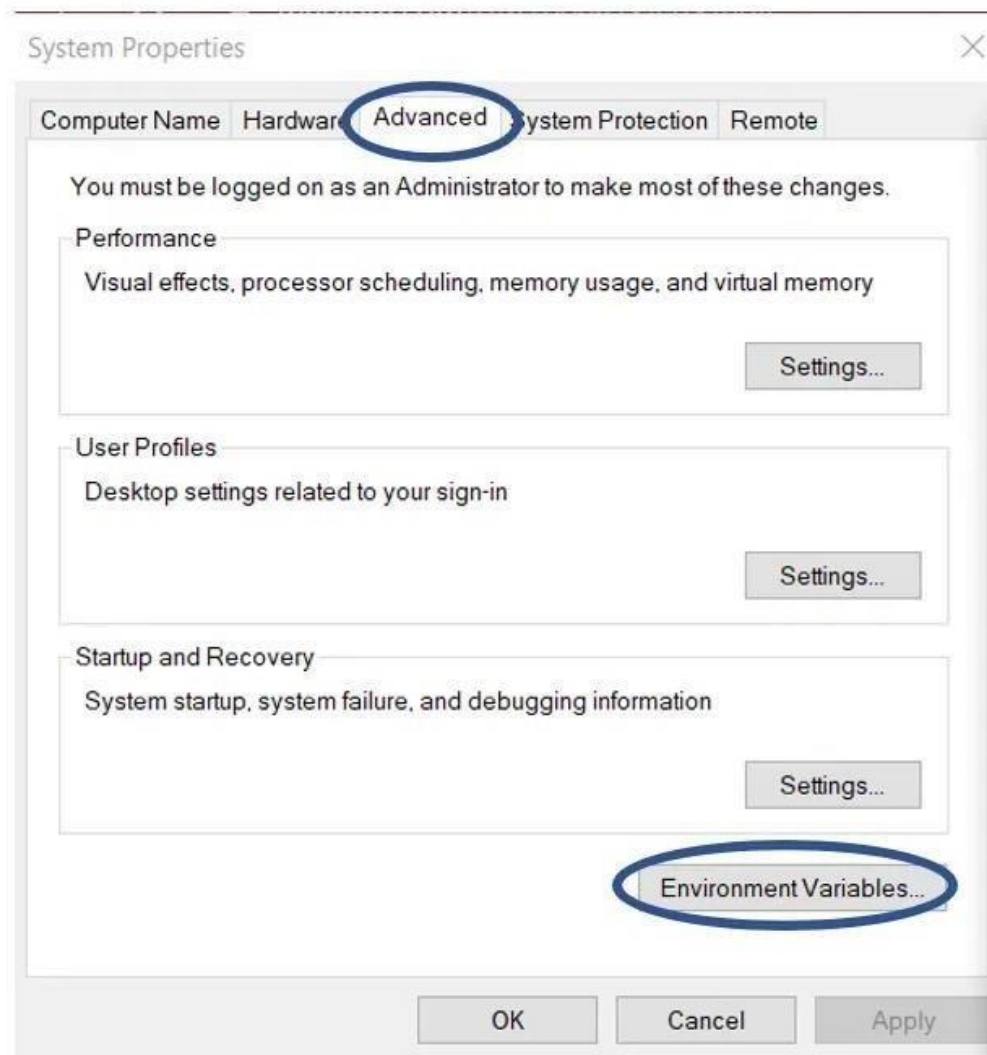


Fig. 6:- Advanced System Settings Screen

3.1 Setting JAVA_HOME

Open environment Variable and click on -New in -User Variable

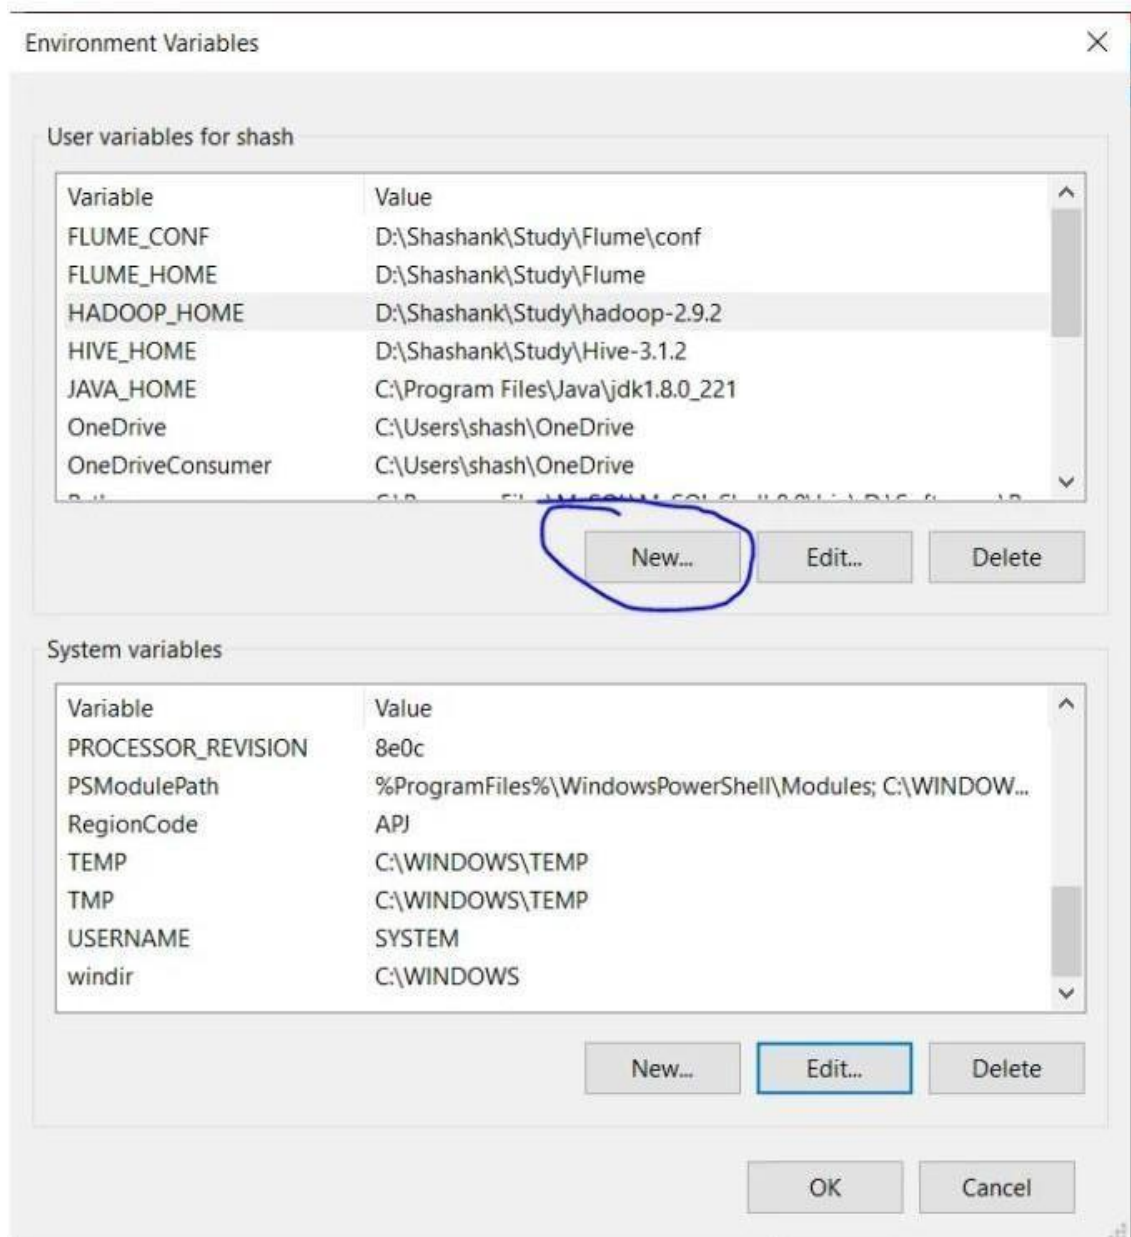


Fig. 7:- Adding Environment Variable

On clicking -Newl, we get below screen.




Fig. 8:- Adding JAVA_HOME

Now as shown, add JAVA_HOME in variable name and path of Java(jdk) in Variable Value. Click OK and we are half done with setting JAVA_HOME.

3.2 Setting HADOOP_HOME

Open environment Variable and click on -New in -User Variable

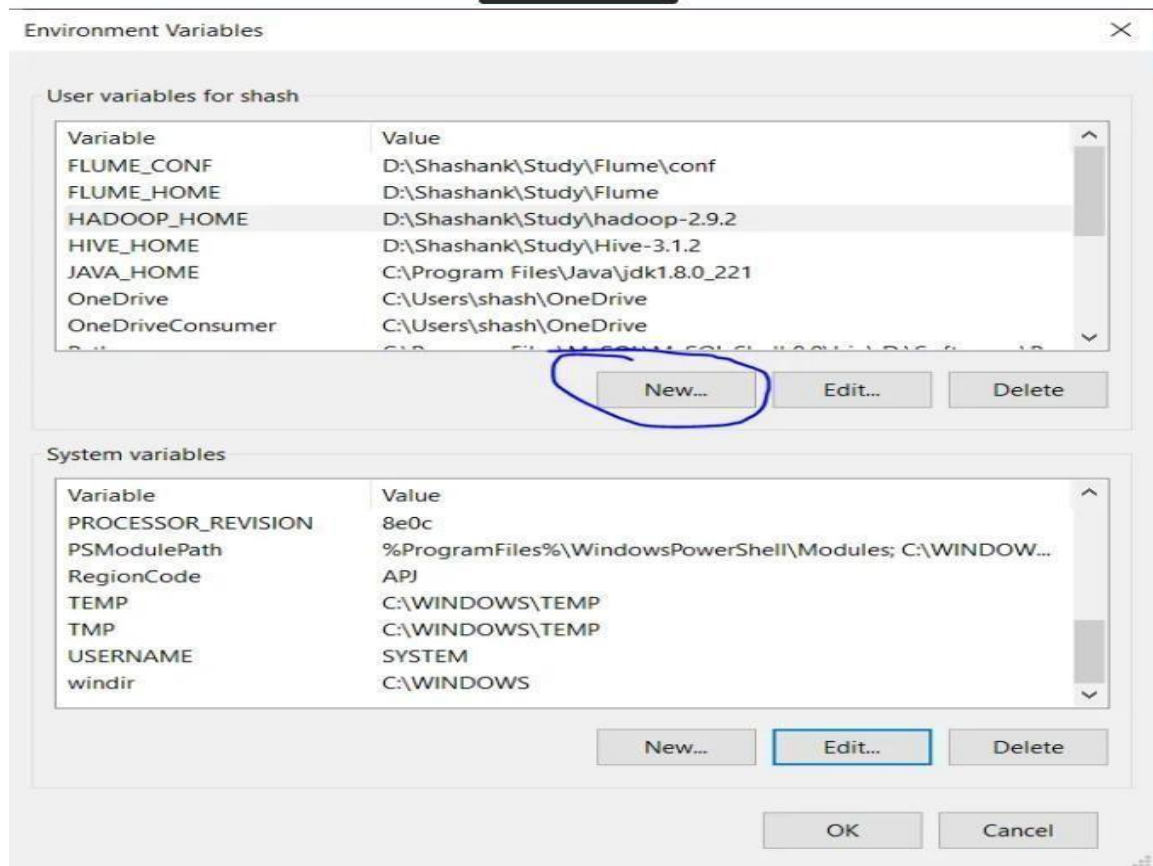


Fig. 9:- Adding Environment Variable

On clicking -New, we get below screen.

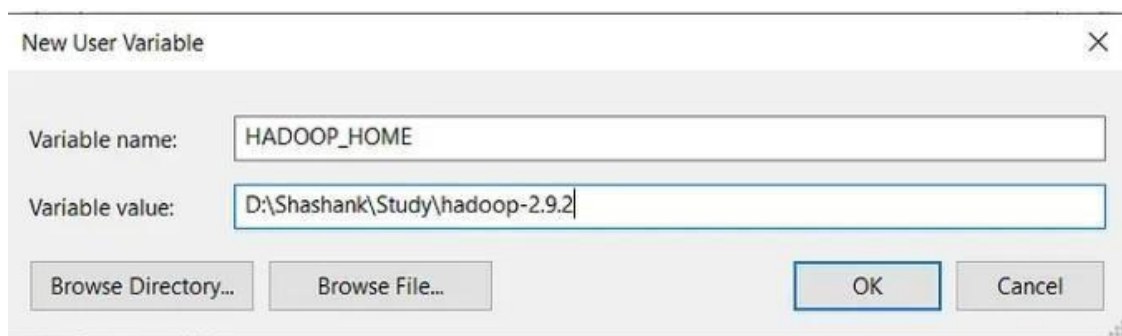


Fig. 10:- Adding HADOOP_HOME

Now as shown, add HADOOP_HOME in variable name and path of Hadoop folder in Variable Value.

Click OK and we are half done with setting HADOOP_HOME.

Note:- If you want the path to be set for all users you need to select -New from System Variables.

3.3 Setting Path Variable

Last step in setting Environment variable is setting Path in System Variable.




Fig. 11:- Setting Path Variable

Select Path variable in the system variables and click on -Editl.

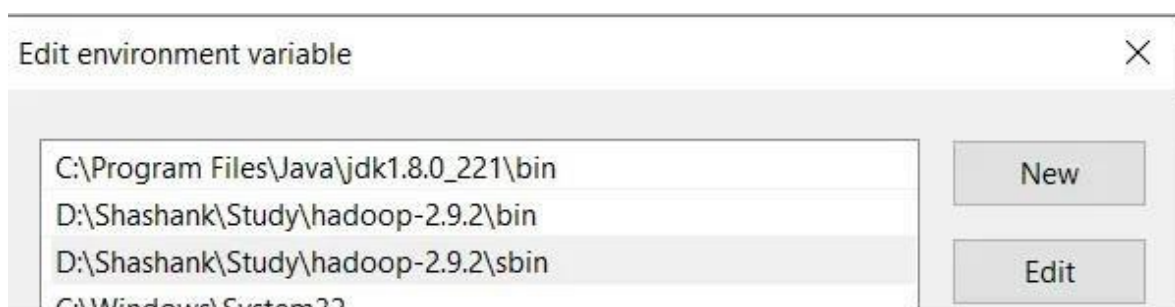


Fig. 12:- Adding Path

Now we need to add these paths to Path Variable one by one:-

- * %JAVA_HOME%\bin
- * %HADOOP_HOME%\bin
- * %HADOOP_HOME%\sbin

Click OK and OK. & we are done with Setting Environment Variables.

3.4 Verify the Paths

Now we need to verify that what we have done is correct and reflecting.

Open a NEW Command Window

Run following commands

```
echo %JAVA_HOME%
echo %HADOOP_HOME%
echo %PATH%
```

4. Editing Hadoop files

Once we have configured the environment variables next step is to configure Hadoop. It has 3 parts:-

4.1 Creating Folders

We need to create a folder data in the hadoop directory, and 2 sub folders namenode and datanode

PC > Shashank (D:) > Shashank > Study > hadoop-2.9.2

Name	Date modified	Type	Size
bin	20-09-2020 16:35	File folder	
data	20-09-2020 16:31	File folder	
etc	13-11-2018 20:45	File folder	
include	13-11-2018 20:45	File folder	
lib	13-11-2018 20:45	File folder	
libexec	13-11-2018 20:45	File folder	
logs	20-09-2020 16:41	File folder	
sbin	13-11-2018 20:45	File folder	
share	13-11-2018 20:45	File folder	
LICENSE	13-11-2018 20:45	TXT File	104 KB
NOTICE	13-11-2018 20:45	TXT File	16 KB
README	13-11-2018 20:45	TXT File	2 KB

Fig. 13:- Creating Data Folder

Create DATA folder in the Hadoop directory

PC > Shashank (D:) > Shashank > Study > hadoop-2.9.2 > data

Name	Date modified	Type	Size
datanode	25-12-2020 15:34	File folder	
namenode	25-12-2020 15:34	File folder	

Fig. 14:- Creating Sub-folders

Once DATA folder is created, we need to create 2 new folders namely, namenode and datanode inside the data folder

These folders are important because files on HDFS resides inside the datanode.

4.2 Editing Configuration Files

Now we need to edit the following config files in hadoop for configuring it :-

(We can find these files in Hadoop -> etc -> hadoop)

- * core-site.xml
- * hdfs-site.xml
- * mapred-site.xml
- * yarn-site.xml
- * hadoop-env.cmd

4.2.1 Editing core-site.xml

Right click on the file, select edit and paste the following content within <configuration> </configuration> tags.

Note:- Below part already has the configuration tag, we need to copy only the part inside it.

```
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
</configuration>
```

4.2.2 Editing hdfs-site.xml

Right click on the file, select edit and paste the following content within

<configuration></configuration>tags.

Note:- Below part already has the configuration tag, we need to copy only the part inside it.

Also replace PATH~1 and PATH~2 with the path of namenode and datanode folder that we created recently(step 4.1).

```
<configuration>
<property>
  <name>dfs.replication</name>
  <value>1</value>
</property>
<property>
  <name>dfs.namenode.name.dir</name>
  <value>C:\hadoop\data\namenode</value>
</property>
<property>
  <name>dfs.datanode.data.dir</name>
  <value>C:\hadoop\data\datanode</value>
</property>
</configuration>
```

4.2.3 Editing mapred-site.xml

Right click on the file, select edit and paste the following content within <configuration>

</configuration> tags.

Note:- Below part already has the configuration tag, we need to copy only the part inside it.

```
<configuration>
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
</configuration>
```

4.2.4 Editing yarn-site.xml

Right click on the file, select edit and paste the following content within <configuration>

</configuration> tags.

Note:- Below part already has the configuration tag, we need to copy only the part inside it.

```
<configuration>
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>
```

4.2.5 Verifying hadoop-env.cmd

Right click on the file, select edit and check if the JAVA_HOME is set correctly or not.

We can replace the JAVA_HOME variable in the file with your actual JAVA_HOME that we configured in the System Variable.

set JAVA_HOME=%JAVA_HOME%

OR

set JAVA_HOME="C:\Program Files\Java\jdk1.8.0_221"

4.3 Replacing bin

Last step in configuring the hadoop is to download and replace the bin folder.

- * Go to this [GitHub Repo](#) and download the bin folder as a zip.
 - * Extract the zip and copy all the files present under bin folder to %HADOOP_HOME%\bin
- Note:- If you are using different version of Hadoop then please search for its respective bin folder and download it.

5. Testing Setup

Congratulation..!!!!

We are done with the setting up the Hadoop in our System.

Now we need to check if everything works smoothly...

5.1 Formatting Namenode

Before starting hadoop we need to format the namenode for this we need to start a NEW Command Prompt and run below command

hadoop namenode -format

[illegible]

Fig. 15:- Formatting Namenode

Note:- This command formats all the data in namenode. So, its advisable to use only at the start and do not use it every time while starting hadoop cluster to avoid data loss.

5.2 Launching Hadoop

Now we need to start a new Command Prompt remember to run it as administrator to avoid permission issues and execute below commands
start-all.cmd

```
C:\Users\shash>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.
starting yarn daemons
'C:\Program' is not recognized as an internal or external command,
operable program or batch file.

C:\Users\shash>
```

Fig. 16:- start-all.cmd

This will open 4 new cmd windows running 4 different Daemons of hadoop:-

- * Namenode
- * Datanode
- * ResourceManager
- * Nodemanager



Fig. 17:- Hadoop Daemons

Note:- We can verify if all the daemons are up and running using jps command in new cmd window.

6. Running Hadoop (Verifying Web UIs)

6.1 Namenode

Open localhost:50070 in a browser tab to verify namenode health.



Fig. 18:- Namenode Web UI

6.2 Resourcemanager

Open localhost:8088 in a browser tab to check resourcemanager details.

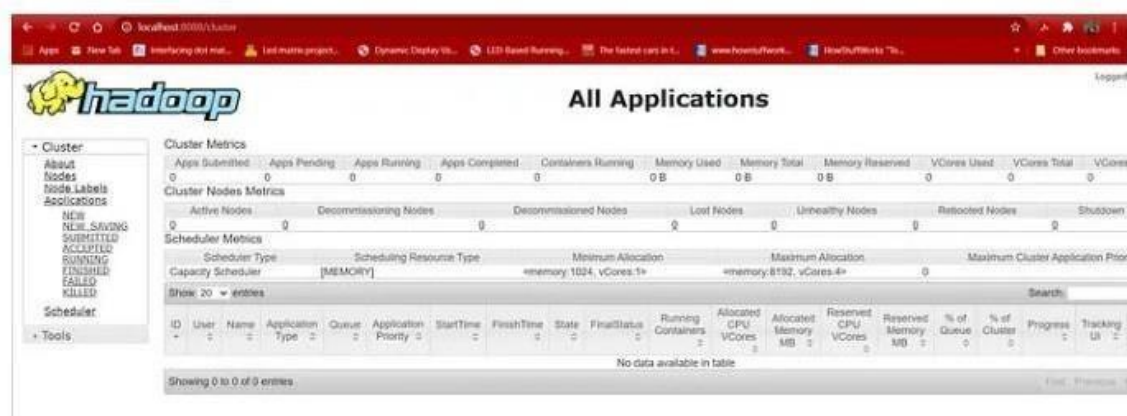


Fig. 19:- Resourcemanager Web UI

6.3 Datanode

Open localhost:50075 in a browser tab to checkout datanode.

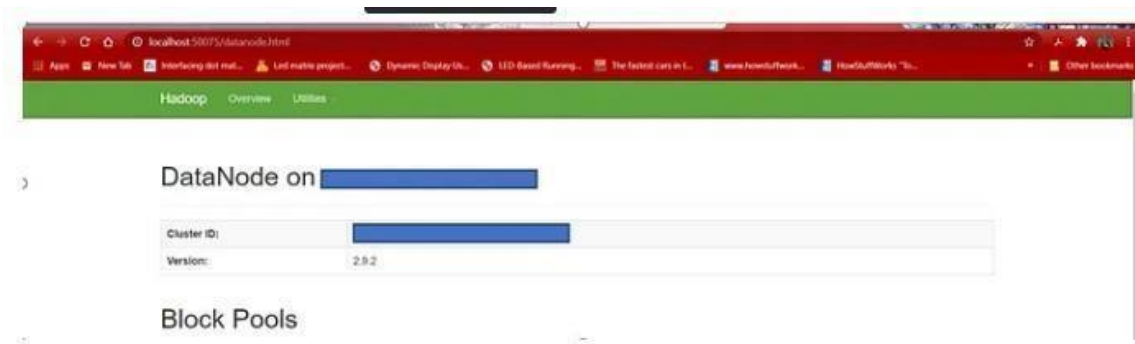


Fig. 20:- Datanode Web UI

OUTPUT:

EXPERIMENT:**Implement Word Count/ Frequency Programs Using Map Reduce.****PROGRAM:**

AIM: To count a given number using map reduce functions.

Hadoop Streaming API for helping us passing data between our Map and Reduce code via STDIN (standard input) and STDOUT (standard output).

Note : Change the file has execution permission (chmod +x /home/hduser/mapper.py)

Change the file has execution permission (chmod +x /home/hduser/reducer.py)

Mapper program

```
mapper.py
import sys
# input comes from STDIN (standard input)
for line in sys.stdin:
    line = line.strip() # remove leading and trailing whitespace
    words = line.split()# split the line into words
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        # tab-delimited; the trivial word count is 1
        print '%s\t%s' % (word, 1)
```

Reducer program

```
"""reducer.py"""
from operator import itemgetter
import sys
current_word = None
current_count = 0
word = None

# input comes from STDIN
for line in sys.stdin:
    line = line.strip() # remove leading and trailing whitespace
    # parse the input we got from mapper.py
    word, count = line.split("\t", 1)
    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue

    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
```

```
if current_word:
    # write result to STDOUT
    print '%s\t%s' % (current_word, current_count)
    current_count = count
    current_word = word
```

do not forget to output the last word if needed!

```
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

Test the code (cat data | map | sort | reduce)

```
hduser@ubuntu:~$ echo "foo foo quux labs foo bar quux" | /home/hduser/mapper.py
```

```
foo 1
foo 1
quux 1
labs 1
foo 1
bar 1
quux 1
```

```
hduser@ubuntu:~$ echo "foo foo quux labs foo bar quux" | /home/hduser/mapper.py | sort -k1,1 | /home/hduser/reducer.py
```

```
bar 1
foo 3
labs 1
quux 2
```

```
hduser@ubuntu:~$ cat /tmp/gutenberg/20417-8.txt | /home/hduser/mapper.py
```

```
The 1
Project 1
Gutenberg 1
EBook 1
of 1
```

OUTPUT: